



G1


Priority queue and syscall read/write

Mikkel Enevoldsen

Kristian Høi

Simon Van Beest

February 19, 2016



A Heap-based priority Queue

1. We have implemented the queue as a max heap. Our datastructure has three items: an `int* root`, `int size` and `int capacity`. The root is an integer pointer to the start of our heap. For navigating in our heap, we have created three macros in `queue.h`: `PARENT`, `LEFT` and `RIGHT`. `size` keeps track of the size of the heap
`int capacity` keeps track of the items we have allocated space to.
In `queue_pop` we simply return the value in the root. Then we set the last item of the queue, the root item, and call `max_heapify` which restores the max heap attribute.
In `queue_push` we insert our new priority at the end of our heap. Then we bottom-up compare the priority with the parent and swap if needed.
If the number of items is equal to the capacity we allocated, we call `realloc`. `realloc` changes the size of the memory block pointer depending on the argument pointer. By this we can add more capacity.
According to the main page `realloc` can fail and return a null pointer. This is why we check if the returned pointer is different from `NULL`. Fortunately `realloc` does not change the original memory block if it fails.
2. `queue_pop` returns 1 if the `queue_pop` is called on an empty queue.
A queue can be initialized again after `queue_destroy` has been called, since we only free the allocated space for the queue's items.
3. We added a test target to our Makefile. By running `make test` it runs the command `python3 bounce.py`
We have tried `bounce.py` with a different amount of tests ranging from 50-1000 and all did pass.

Kudos System Calls for Basic I/O

1. We have created the file `rw.c` which contains various systemcalls to read and write. This has been placed in userland, as it is a testfile and not part of the kernel.
To handle these calls from userland we created `read.c`, `read.h`, `write.c` and `write.h`. These are placed in `kudos/proc` as the documentation of kudos suggests.
2. The files mentioned above are made following the hello syscall example presented at the exercises. The fact that the userland program and the kernel are separated suggests that the implementation is bullet proof.
3. In the file, `rw.c`, we use `syscall` to access kernel mode for the I/O processes. The testing is based on visual feedback. We call `syscall_write` on a predefined string, and reads the user input with `syscall_read` and write it with `syscall_write`. If the test is successful it will print the user's input in the terminal. It is only possible to enter one character (or paste a string) at a time. We call `syscall_halt` to avoid kernel panic.